Sequential Testing for A/B Experiments

A Focused Guide in Three Acts

Introduction

The problem. Standard hypothesis tests guarantee a controlled false positive rate only when results are analysed once, at a pre-specified sample size. In practice, teams routinely monitor experiments and peek at interim results — inflating the Type I error rate well beyond the nominal α.

The solution. *Sequential testing* provides statistical methods whose error guarantees hold regardless of the number or timing of analyses. The confidence intervals become wider, but it can still reduce overall runtime if effects become clear at an early stage.

This guide. Three acts, each with an interactive simulation, technical exposition, and key formulas:

1. **The Peeking Problem** — Why interim analyses invalidate standard tests.
2. **The Eppo Solution**— Confidence sequences, variance reduction, and the recommended **hybrid approach**.
3. **Alternative Methods**— Bonferroni, Pocock, and O'Brien–Fleming corrections for teams without a sequential testing platform.

Target audience. Data scientists and analysts familiar with A/B testing, hypothesis testing, and confidence intervals.

Act 1 — The Peeking Problem

Simulation

Below is a simulation of an A/B test. Specify the effect size, number of users, and significance level. The plot shows the running estimate of the relative uplift and the standard 95% confidence interval.

Simulation 1: fixed-horizon confidence intervals.

Relative effect size (%): 0%

Control conversion rate (10.0%)

Number of users (n = 10000)

α (0.05)

Default: α = 0.05

Relative effect size (%) (0%)

Default: effect size = 0%

Number of peeks (K) (14)

Estimated Power (1 - β)

Inapplicable since there is no true effect

Computed from effect size, n, and α (two-sided z-test approximation). Shown as inapplicable when effect size is 0.

Re-randomize

Standard 95% CI

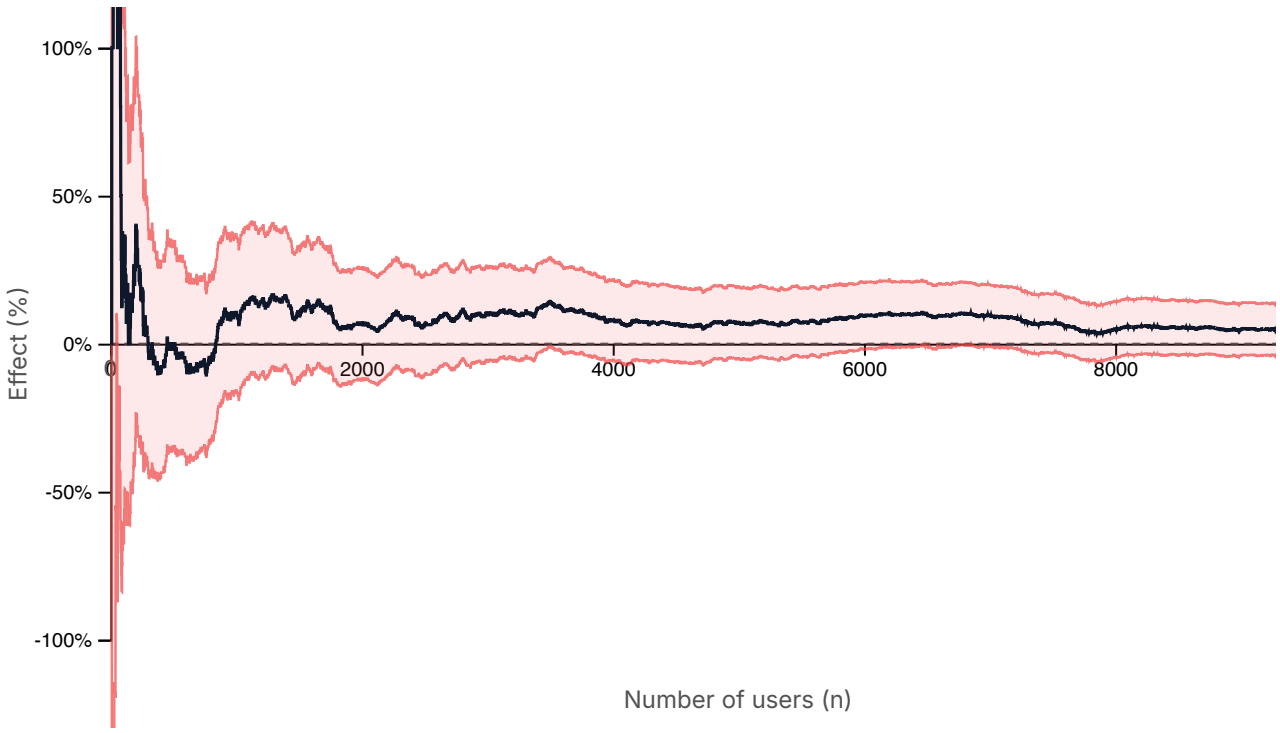

In this simulation, would peeking 14 times at equal time intervals during the test show at least one statistically significant result?

No

Share of simulations in which peeking 14 times at equal time intervals during the test would show at least one statistically significant result, across 1000 repetitions:

Run 1000 repetitions

Click the button to run simulations

Hide simulation assumptions and notes

Assumptions

- Binary outcome metric (Bernoulli), modeled as conversion in each arm.
- Uplift is modeled as relative lift on baseline: $\text{uplift} = (\bar{p}_B - \bar{p}_A)/\bar{p}_A$, with generator $p_B = p_A(1 + \text{lift})$.
- Null hypothesis is $H_0 : \text{lift} = 0$ (equivalently, no treatment effect).
- Standard deviation follows Bernoulli variance in each arm: $\sigma_A = \sqrt{p_A(1 - p_A)},\ \sigma_B = \sqrt{p_B(1 - p_B)}$, so $\text{SE}(\hat{p}_B - \hat{p}_A) = \sqrt{(\sigma_A^2 + \sigma_B^2)/n}$ where n is the per-group sample size.
- Equal traffic split: 50% control and 50% treatment.
- Control conversion baseline is user-specified via slider.
- Independent users/events within and across arms (no clustering or interference).
- No missing data, no delayed outcomes, and no sample-ratio mismatch.
- Two-sided significance check at each look: confidence interval crossing zero is treated as significant.
- Peeks occur at exactly K equal time intervals over the test duration.

Notes

- The Monte Carlo estimate uses 1000 repetitions, so displayed rates include simulation noise.
- Example: baseline 10% with +4% relative lift gives treatment 10.4%.
- Estimated power is shown as $\approx 100\%$ when extremely close to 1, rather than exactly 100%.

Why does this happen?

Standard confidence intervals are valid only when evaluated once at a pre-specified sample size. Each additional peek gives another chance to get a "significant" result by luck.

How bad is it?

Values below are calibrated to the default simulation settings ($\alpha = 0.05$, $n = 10000$) using equal-interval peeks.

Checking schedule	Actual false positive rate
Once at the end (as designed)	$\sim 5\%$
Daily for 1 week	$\sim 17\%$
Daily for 2 weeks	$\sim 22\%$
Daily for 4 weeks (28 looks)	$\sim 29\%$
Continuously (every observation)	Can exceed 60%

With daily checks over a 2-week experiment, about one in five “significant” results will be a false positive. Features released on this basis may have no real effect.

Key Takeaway

Sequential methods control the probability of ever making a false positive, no matter when we stop. They use stricter or time-dependent thresholds to account for repeated checking. This way, we can monitor results as often as we like, and the error rate stays controlled.

Hide the math

The math

Why peeking inflates Type I error

The simulation tracks the **relative uplift**: the treatment effect expressed as a percentage of the control mean. Formally:

$$\hat{u}_n = 100 \cdot \frac{\bar{X}_{B,n} - \bar{X}_{A,n}}{\bar{X}_{A,n}}$$

- $\hat{u}_n$ — estimated relative uplift (in %) after n observations per group
- $\bar{X}_{B,n}$ — running mean outcome in the **treatment** group after n observations
- $\bar{X}_{A,n}$ — running mean outcome in the **control** group after n observations
- n — number of observations (e.g. users) in each group so far

We test $H_0 : u = 0$ (no relative uplift). A standard 95% confidence interval for the uplift, valid at one pre-specified look, is:

$$\hat{u}_n \pm 100 \cdot \frac{\widehat{\mathrm{SE}}_n}{\bar{X}_{A,n}} \cdot 1.96$$

- 1.96 — the critical value for a 95% confidence interval (the 97.5th percentile of the standard Normal distribution)
- $\widehat{\mathrm{SE}}_n$ — estimated standard error of the *absolute* difference in means $(\bar{X}_{B,n} - \bar{X}_{A,n})$, equal to $\sqrt{(\hat{\sigma}_A^2 + \hat{\sigma}_B^2)/n}$; dividing by $\bar{X}_{A,n}$ converts this to the SE of the relative uplift (delta method)
- $\hat{\sigma}_A^2$, $\hat{\sigma}_B^2$ — estimated variance of outcomes in the control and treatment groups

The false positive rate at a single planned look is:

$$\Pr(\text{reject } H_0 \text{ at one look}) = \alpha = 0.05$$

- α — the significance level; the maximum acceptable probability of a false positive (here 5%)

If we repeatedly check the data, we effectively run many tests. If the looks were independent, the probability of at least one false positive in K looks would be:

$$\Pr(\text{at least one false positive in } K \text{ looks}) = 1 - (1 - \alpha)^K$$

- K — total number of times we look at (peek at) the data during the experiment

A useful approximation for small α:

$$1 - (1 - \alpha)^K \approx K\alpha$$

Real interim looks are positively correlated, so the exact value is lower than the independence formula, but still substantially above α. That is why repeated peeking inflates Type I error.

Act 2 — The Eppo Solution

Simulation

This is the same simulation as Act 1, plus Eppo's sequential confidence interval (blue) on top of the standard fixed-horizon confidence interval (red). The blue interval is typically wider at any single look, but it keeps error control valid under repeated monitoring.

Simulation 2: fixed-horizon + Eppo sequential confidence intervals.

Relative effect size (%): 0%

Control conversion rate (10.0%)

Number of users (n = 10000)

α (0.05)

Default: α = 0.05

Relative effect size (%) (0%)

Default: effect size = 0%

Number of peeks (K) (14)

Estimated Power (1 - β)

Inapplicable since there is no true effect

Computed from effect size, n, and α (two-sided z-test approximation). Shown as inapplicable when effect size is 0.

Re-randomize

Standard 95% CI Sequential CI (Eppo)

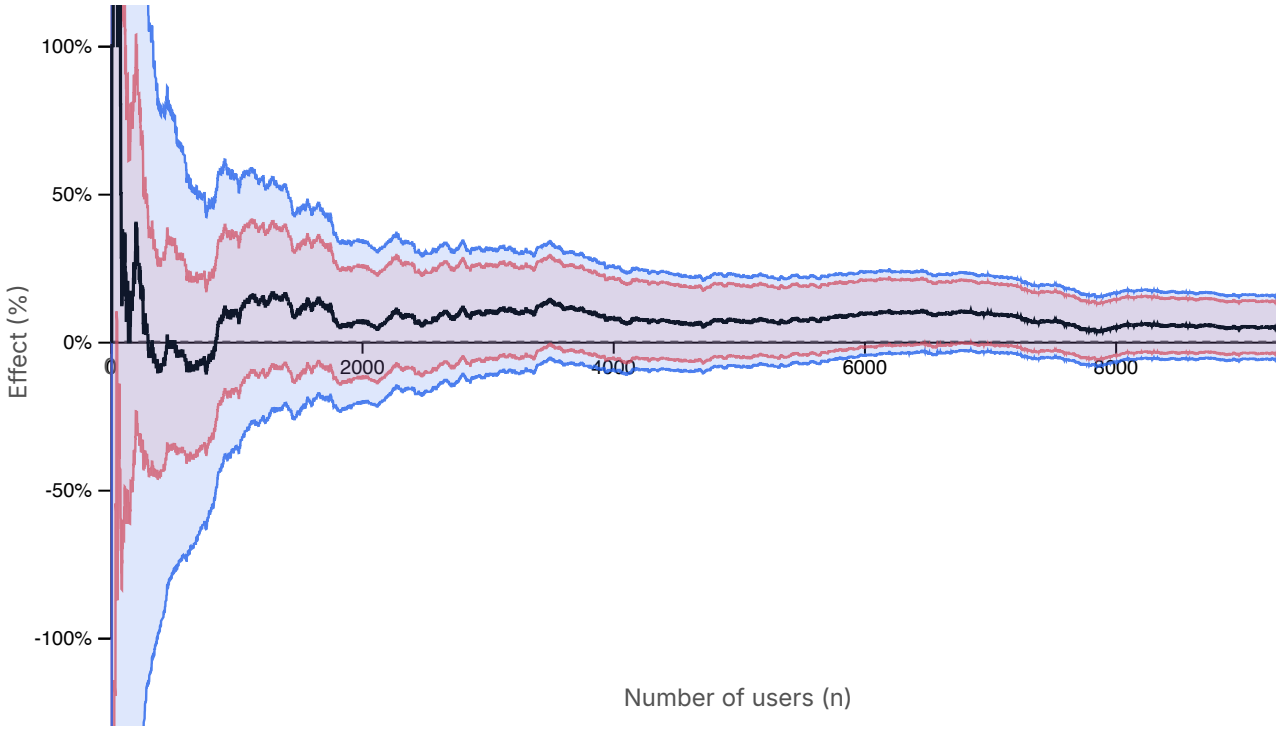

In this simulation, would peeking 14 times at equal time intervals during the test show at least one statistically significant result?

No

Share of simulations in which peeking 14 times at equal time intervals during the test would show at least one statistically significant result, across 1000 repetitions:

Run 1000 repetitions

Method	Share crossing	Mean |effect| when significant	Mean |effect| at end of test when significant
Standard 95% CI	22.5%	21.6%	5.4%
Sequential CI (Eppo)	5.9%	17.1%	8.2%

Hide simulation assumptions and notes

Assumptions

- Binary outcome metric (Bernoulli), modeled as conversion in each arm.
- Uplift is modeled as relative lift on baseline: $\text{uplift} = (\bar{p}_B - \bar{p}_A)/\bar{p}_A$, with generator $p_B = p_A(1 + \text{lift})$.
- Null hypothesis is $H_0 : \text{lift} = 0$ (equivalently, no treatment effect).
- Standard deviation follows Bernoulli variance in each arm: $\sigma_A = \sqrt{p_A(1 - p_A)},\ \sigma_B = \sqrt{p_B(1 - p_B)}$, so $\text{SE}(\hat{p}_B - \hat{p}_A) = \sqrt{(\sigma_A^2 + \sigma_B^2)/n}$ where n is the per-group sample size.
- Equal traffic split: 50% control and 50% treatment.
- Control conversion baseline is user-specified via slider.
- Independent users/events within and across arms (no clustering or interference).
- No missing data, no delayed outcomes, and no sample-ratio mismatch.
- Two-sided significance check at each look: confidence interval crossing zero is treated as significant.
- Peeks occur at exactly K equal time intervals over the test duration.

Notes

- The Monte Carlo estimate uses 1000 repetitions, so displayed rates include simulation noise.
- Example: baseline 10% with +4% relative lift gives treatment 10.4%.
- Estimated power is shown as $\approx 100\%$ when extremely close to 1, rather than exactly 100%.

Why does this work?

The fixed-horizon confidence interval controls error for one planned analysis. Eppo's sequential confidence interval uses a time-dependent boundary, so the guarantee is valid no matter when or how often you peek.

Tradeoff: at any single look, the interval is wider; benefit: no hidden inflation from repeated checks.

How different is it?

Values below are calibrated to default settings ($\alpha = 0.05$, $n = 10000$, baseline 10%, no true effect) using equal-interval peeks.

Checking schedule	Standard 95% confidence interval	Sequential confidence interval (Eppo)
Once at the end (1 look)	$\sim 5\%$	$\sim 2\%$
Daily for 1 week (7 looks)	$\sim 16\%$	$\sim 4\%$
Daily for 2 weeks (14 looks)	$\sim 22\%$	$\sim 5\%$
Daily for 4 weeks (28 looks)	$\sim 28\%$	$\sim 5\%$
Very frequent checks (50 looks)	$\sim 33\%$	$\sim 6\%$

Under repeated peeking, the standard confidence interval inflates false positives substantially, while the sequential confidence interval stays close to the target error level.

Note: the sequential CI shows ~2% at a single look — below the nominal 5% — because the sequential multiplier is calibrated for continuous monitoring and is intentionally conservative at any fixed look.

Key Takeaway

Sequential intervals trade narrowness at a single look for validity across all looks.

your team monitors live, this is the right tradeoff.

The Hybrid Approach

The sequential confidence interval is always wider than a fixed-horizon confidence interval at any given sample size. For the **primary KPI** — the metric the experiment is designed to move — this width penalty reduces statistical power.

Apply sequential testing where early stopping adds the most value (safety guardrails), and fixed-horizon testing where statistical power matters most (the primary metric).

The hybrid approach partitions metrics into two categories:

- **Guardrail KPIs** (revenue, error rate, latency, etc.): monitored continuously with a **sequential confidence interval**. If the confidence interval excludes zero on the harmful side, abort the experiment immediately.
- **Primary KPI** (the metric the experiment targets): analysed with a **standard fixed-horizon confidence interval** at the pre-planned end date. No loss of power (no sequential correction applied).

What you gain

	Full sequential	Hybrid
Primary KPI confidence interval width	Wider (~10–40%)	Standard (no penalty)
Primary KPI power	Reduced	Full
Guardrail protection	Continuous	Continuous
Early stopping for success	Yes	No
Early stopping for harm	Yes	Yes
Complexity	Higher	Lower

Key Takeaway

The hybrid approach in one sentence: Monitor guardrail KPIs with a sequential confidence interval for early abort; analyse the primary KPI with a standard confidence interval at the planned end-date.

This gives you the safety of sequential testing where it matters most (preventing harm) without paying the power penalty on the metric you care about most (the primary KPI).

This is the recommended approach for most A/B testing programmes.

Hide the math

The math

1. Fixed-horizon confidence interval (from Act 1)

A standard 95% confidence interval for the relative uplift $\hat{u}_n$, valid at one pre-specified sample size n, is:

$$\hat{u}_n \pm 100 \cdot \frac{\widehat{\mathrm{SE}}_n}{\bar{X}_{A,n}} \cdot 1.96$$

- $\hat{u}_n$ — estimated relative uplift (in %) after n users per group have been observed
- 1.96 — the critical value for a 95% confidence interval (the 97.5th percentile of the standard Normal distribution)
- $\widehat{\mathrm{SE}}_n$ — estimated standard error of the difference in means, equal to $\sqrt{\hat{\sigma}_A^2/n + \hat{\sigma}_B^2/n}$
- $\hat{\sigma}_A^2$ — estimated variance of outcomes in the control group (group A)
- $\hat{\sigma}_B^2$ — estimated variance of outcomes in the treatment group (group B)
- $\bar{X}_{A,n}$ — running mean outcome in the control group after n observations
- n — number of users assigned to each group so far

2. Sequential confidence interval (Eppo)

Eppo's sequential confidence interval replaces the fixed multiplier 1.96 with one that depends on the number of observations:

$$\hat{u}_n \pm 100 \cdot \frac{\widehat{\mathrm{SE}}_n}{\bar{X}_{A,n}} \sqrt{\frac{n+\nu}{n} \log\left(\frac{n+\nu}{\nu\,\alpha}\right)}$$

- $\hat{u}_n$ — estimated relative uplift (in %) after n users per group
- $\widehat{\mathrm{SE}}_n$ — estimated standard error of the difference in means at the current sample size
- $\bar{X}_{A,n}$ — running mean outcome in the control group (used to convert absolute SE to relative %)
- n — current number of users in each group
- ν — tuning parameter that controls the width–power tradeoff, calibrated to the planned sample size. See the formula below.
- α — target Type I error level (e.g. 0.05 for a 5% false-positive rate)
- $\log$ — natural logarithm

3. The multiplier

The key difference between the two formulas is the multiplier of the standard error. In the fixed-horizon case it is the constant 1.96. In Eppo's sequential implementation it is:

$$m(n) = \sqrt{\frac{n+\nu}{n} \log\left(\frac{n+\nu}{\nu\,\alpha}\right)}$$

- $m(n)$ — the time-varying multiplier applied to the standard error (replaces 1.96 from the fixed-horizon formula)
- n — current number of users in each group
- ν — tuning parameter (explained below)
- α — target Type I error level (e.g. 0.05)
- $\log$ — natural logarithm

This multiplier is larger than 1.96, which is what keeps the Type I error controlled under continuous monitoring. The multiplier is especially high when n is small. As n grows, $m(n)$ first decreases — so the confidence interval narrows — reaches a minimum somewhere before the planned sample size n^*, and then slowly increases again. The interval is therefore tightest in the middle of the experiment and widens slightly if the experiment runs well past n^*. It always remains above 1.96.

The tuning parameter ν controls this trade-off between early-stopping power and long-run width. Eppo sets it as:

$$\nu = \frac{n^*}{\log(n^*/\alpha) - 1}$$

- ν — tuning parameter that balances the interval width at early looks versus the planned end date
- n^* — the planned (maximum) sample size per group — the horizon at which the experiment is expected to end
- α — target Type I error level

Setting ν too high (relative to actual traffic) shifts the minimum of $m(n)$ to a sample size you never reach, so the interval stays wider than necessary throughout the experiment. Setting it too low shifts the minimum to a point you pass quickly, after which the multiplier starts rising earlier than optimal. The formula above calibrates ν so the multiplier is near its minimum around n^* — the point where you expect to make a decision. Being off by a factor of two is fine; being off by a factor of ten matters.

4. Time-uniform coverage guarantee

The sequential CI satisfies:

$$\Pr(u \in \mathrm{CI}_n \text{ for all } n \geq 1) \geq 1 - \alpha$$

- u — the true (unknown) relative uplift of the treatment compared to control
- CI_n — the sequential confidence interval computed at sample size n
- $\Pr(\cdot)$ — probability over repeated experiments
- α — target Type I error level; the guarantee holds with probability at least $1 - \alpha$

So unlike fixed-horizon confidence intervals, the guarantee still holds under continuous monitoring.

As n increases, the multiplier $m(n)$ first decreases and then slowly rises — so the confidence interval narrows through most of the experiment and widens slightly if the experiment runs past the planned sample size.

Act 3 — Alternative Methods

Three group sequential methods for controlling false positives under interim analyses, for teams implementing sequential monitoring without a dedicated platform.

How these methods differ

Suppose you plan to analyse your experiment at K pre-specified interim time points. Using the standard $\alpha = 0.05$ threshold at each analysis inflates the overall Type I error rate (Act 1).

The fix: **adjust the per-analysis significance level** so that the family-wise error rate remains at α. The three methods below differ in *how* they allocate the error budget across the K analyses.

- **Bonferroni:** equal allocation, α/K per analysis. Conservative because it ignores correlation across analyses.
- **Pocock:** constant critical value calibrated to the joint distribution of the test statistics across analyses. Tighter than Bonferroni.
- **O'Brien–Fleming:** front-loaded allocation. Very strict early, nearly standard at the final analysis.

Eppo's approach (Act 2) does not require pre-specifying K— it provides a continuously valid guarantee.

Simulation

This extends the Act 1/2 simulation by adding Bonferroni, Pocock, and O'Brien–Fleming confidence intervals, so you can compare all methods under the same settings.

Method 1: Bonferroni

The simplest correction. Divide α by K. One line of code.

Method 1: Bonferroni Correction (Simplest)

The idea

Before the experiment starts, fix the number of planned analyses K (e.g. weekly for 4 weeks: $K = 4$). At each analysis, use a significance level of α/K instead of α.

The recipe

1. Choose K (number of planned peeks) and α (e.g. 0.05).

2. At each peek $k = 1, \ldots, K$, compute the standard CI using the adjusted significance level:

$$\mathrm{CI}_k = \hat{\tau}(t_k) \ \pm \ \hat{\sigma}_{\hat{\tau}}(t_k) \cdot z_{\alpha/(2K)}$$

where $\hat{\tau}(t_k)$ is the estimated treatment effect at peek k, $\hat{\sigma}_{\hat{\tau}}(t_k)$ is its standard error, and $z_{\alpha/(2K)}$ is the critical value from a Normal distribution at the adjusted significance level.

3. If the CI does not cross 0: stop and declare significance.
4. If $k = K$ and the CI crosses 0: no significant effect.

Critical Values

K (peeks)	α/K	$z_{\alpha/(2K)}$	CI multiplier vs. 1.96
1	0.050	1.96	1.00× (no correction)
2	0.025	2.24	1.14×
4	0.0125	2.50	1.28×
8	0.00625	2.73	1.39×
12	0.00417	2.87	1.46×
20	0.0025	3.02	1.54×
52	0.00096	3.29	1.68× (weekly for a year)

With $K = 4$, replace $z = 1.96$ with $z = 2.50$. The CI is 28% wider at each analysis, but the family-wise error rate is controlled at $alpha$.

Implementation step by step

Given data arrays `control[]` and `treatment[]` at each scheduled peek:

1. **Compute group statistics:**

$$\bar{Y}_0 = \frac{1}{n_0} \sum_{i=1}^{n_0} Y_{0,i}, \qquad s_0^2 = \frac{1}{n_0 - 1} \sum_{i=1}^{n_0} (Y_{0,i} - \bar{Y}_0)^2$$

2. **Compute the estimated treatment effect and its standard error:**

$$\hat{\tau} = \bar{Y}_1 - \bar{Y}_0, \qquad \mathrm{SE} = \sqrt{\frac{s_0^2}{n_0} + \frac{s_1^2}{n_1}}$$

3. **Look up the adjusted critical value:**

$$z^* = z_{\alpha/(2K)} = \Phi^{-1}\left(1 - \tfrac{\alpha}{2K}\right)$$

where Φ^{-1} is the inverse of the standard Normal CDF (e.g. `scipy.stats.norm.ppf` in Python).

4. **Construct the CI and decide:**

$$\mathrm{CI} = \left[\hat{\tau} - z^* \cdot \mathrm{SE},\ \hat{\tau} + z^* \cdot \mathrm{SE}\right]$$

If the CI does not contain 0, stop and declare significance. Otherwise, wait until the next scheduled peek.

Pros and cons

+ Dead simple. One line of code: change 1.96 to $z_{\alpha/(2K)}$.
+ Always valid. Works for any test statistic, any distribution.

– Conservative. The CI is wider than necessary because Bonferroni ignores the correlation between test statistics at successive peeks.
– Scales poorly. As K grows, the correction becomes increasingly harsh.

Method 2: Pocock (1977)

Constant critical boundary. Tighter than Bonferroni by exploiting correlation.

Method 2: Pocock Boundaries (1977)

The idea

Like Bonferroni, pre-specify K equally-spaced analysis times. But instead of Bonferroni's crude α/K split, use a constant critical value c_P that accounts

for the correlation between test statistics at different peeks.

Because the test statistics $Z_1, Z_2, \ldots, Z_K$ at successive analyses share overlapping data, they are positively correlated. Pocock's method exploits this to derive a tighter threshold.

The recipe

1. Choose K and α.
2. Look up (or compute) the Pocock critical value $c_P(K, \alpha)$.
3. At each peek k, compute the z-statistic: $Z_k = \hat{\tau}(t_k)/\hat{\sigma}_{\hat{\tau}}(t_k)$.
4. If $|Z_k| > c_P$: stop and declare significance.

Pocock vs Bonferroni critical values

K	c_P **(Pocock)**	$z_{\alpha/(2K)}$ **(Bonf.)**	**Improvement**
2	2.18	2.24	3% tighter
4	2.36	2.50	6% tighter
8	2.51	2.73	8% tighter
12	2.60	2.87	9% tighter
20	2.69	3.02	11% tighter

Pocock gives a constant threshold that is a few percent lower than Bonferroni's. The key advantage: the threshold is the same at every peek, so it is almost as simple to use as Bonferroni.

How the boundary is computed

Under the null, the z-statistics at successive analyses follow a multivariate Normal with correlation:

$$\mathrm{Cor}(Z_j, Z_k) = \sqrt{\frac{\min(j,k)}{\max(j,k)}} \quad \text{for } 1 \le j, k \le K$$

The Pocock boundary c_P is the unique value satisfying:

$$\mathbb{P}\Big(\max_{1\le k\le K} |Z_k| > c_P\Big) = \alpha$$

Once c_P is known, the decision rule at each peek is identical to Bonferroni, but with c_P replacing $z_{\alpha/(2K)}$.

Method 3: O'Brien–Fleming (1979)

Decreasing threshold. Nearly no penalty at the final analysis.

Method 3: O'Brien–Fleming Boundaries (1979)

The idea

Instead of a constant threshold, use a threshold that *decreases* over time. At the first peek, demand very strong evidence; at the last peek, demand nearly the same evidence as a classical test.

The recipe

At peek k (of K total), use the critical value:

$$c_k^{\text{OBF}} \approx z_{\alpha/2} \cdot \sqrt{K/k}$$

where $z_{\alpha/2} \approx 1.96$ is the standard Normal critical value.

Critical values at each peek (K=4)

Peek k (of K=4)	c_k^{OBF}	c_P (Pocock)	$z_{\alpha/(2K)}$ (Bonf.)	$z_{\alpha/2}$ (classi
1 (25% of data)	**4.05**	2.36	2.50	1.96
2 (50% of data)	**2.86**	2.36	2.50	1.96
3 (75% of data)	**2.34**	2.36	2.50	1.96
4 (100% of data)	**2.02**	2.36	2.50	1.96

O'Brien–Fleming is extremely conservative early (threshold 4.05 at the

first analysis — early stopping is rare) but barely penalises the final analysis (2.02 vs. 1.96 — only 3% wider). This makes it highly efficient when the experiment runs to completion, since the final analysis is only minimally penalized.

Implementation step by step

1. **Before the experiment:** Fix K and α. At each peek k, the fraction of total data collected is $t_k = k/K$ (assuming equal spacing).

2. **At peek k, compute the critical value:**

$$c_k^{\text{OBF}} = \frac{z_{\alpha/2}}{\sqrt{t_k}} = z_{\alpha/2} \cdot \sqrt{\frac{K}{k}}$$

3. **Compute the z-statistic** (same as Bonferroni and Pocock):

$$Z_k = \frac{\hat{\tau}(t_k)}{\text{SE}(t_k)}$$

4. **Decide:** If $|Z_k| > c_k^{\text{OBF}}$: reject. Otherwise: continue to peek $k+1$.

5. **Convert to a confidence interval** (optional but recommended):

$$\text{CI}_k = \hat{\tau}(t_k) \ \pm \ c_k^{\text{OBF}} \cdot \text{SE}(t_k)$$

Note: the CI width *changes* across peeks because c_k^{OBF} varies.

Practical tips

The exact O'Brien–Fleming boundaries (computed via multivariate Normal integration) differ slightly from the $z_{\alpha/2}/\sqrt{t_k}$ approximation. The approximation is accurate for $K \leq 10$ and is what most software implements. For larger K, use a library (e.g. `statsmodels` or the R package `gsDesign`).

How methods compare in simulation

Simulation 3: fixed-horizon, Eppo, and three alternative sequential methods.

Relative effect size (%): 0%

Control conversion rate (10.0%)

Number of users (n = 10000)

α (0.05)

Default: α = 0.05

Relative effect size (%) (0%)

Default: effect size = 0%

Number of peeks (K) (6)

Estimated Power (1 - β)

Inapplicable since there is no true effect

Computed from effect size, n, and α (two-sided z-test approximation). Shown as inapplicable when effect size is 0.

Re-randomize

Standard 95% CI Sequential CI (Eppo) Pocock (K=6) O'Brien–Fleming (K=6) Bonferroni (K=6)

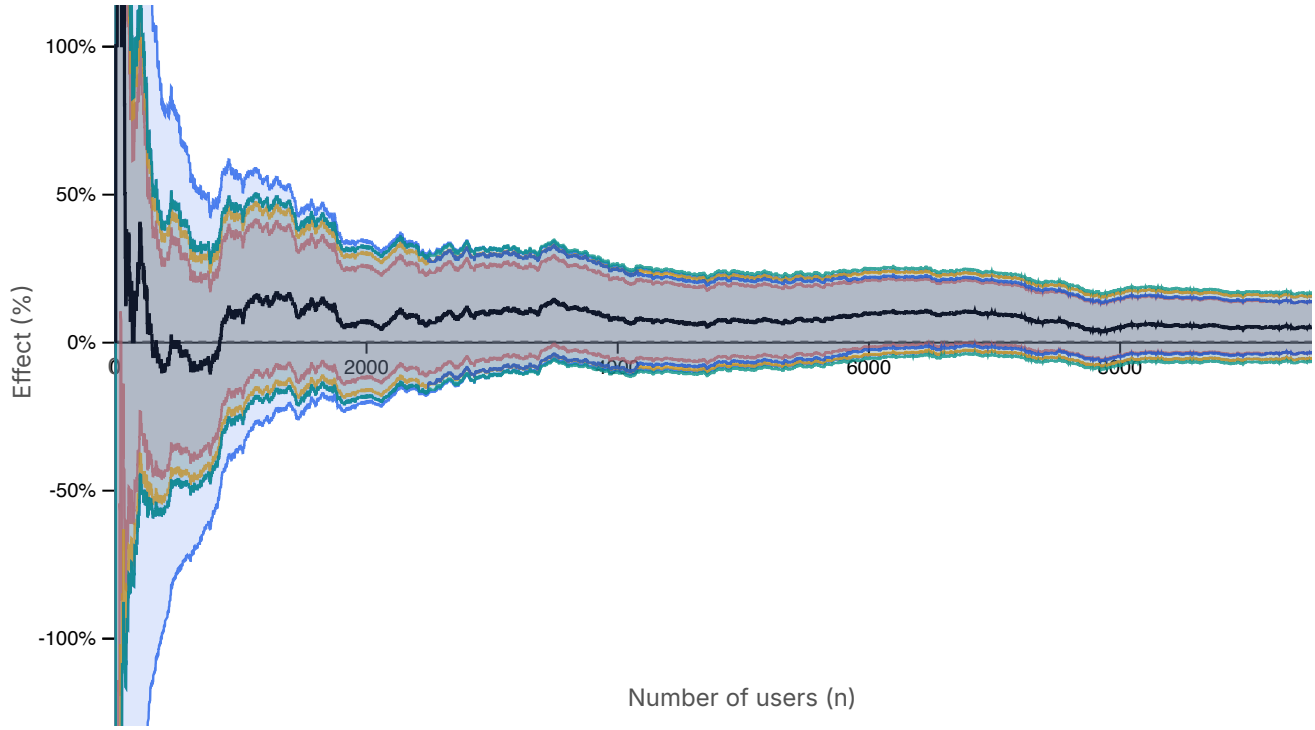

In this simulation, would peeking 6 times at equal time intervals during the test show at least one statistically significant result?

No

Result interpretation: click “Run 1000 repetitions” to estimate how often each method crosses

significance under the current settings.

Bonferroni: most conservative among the three DIY methods (lowest crossing share).
Pocock: less conservative than Bonferroni with the same threshold at each look.
O'Brien–Fleming: very strict early, then close to classical thresholds at later looks.
Sequential confidence interval (Eppo): anytime-valid and typically conservative in this setup.

Share of simulations in which peeking 6 times at equal time intervals during the test would show at least one statistically significant result, across 1000 repetitions:

Run 1000 repetitions

Click the button to run simulations

Show simulation assumptions and notes

Comparison

The confidence interval width rows below are illustrative values for $K = 4$, independent of the slider above.

	Bonferroni	Pocock	OBF	Eppo
Pre-specify K?	Yes	Yes	Yes	No
Threshold type	Constant	Constant	Decreasing	Continuous
Confidence interval at peek 1 (K=4)	$2.50 \times \hat{\sigma}$	$2.36 \times \hat{\sigma}$	$4.05 \times \hat{\sigma}$	$\sim 2.7 \times \hat{\sigma}$
Confidence interval at final analysis (K=4)	$2.50 \times \hat{\sigma}$	$2.36 \times \hat{\sigma}$	$2.02 \times \hat{\sigma}$	$\sim 2.4 \times \hat{\sigma}$
Valid between peeks?	No	No	No	Yes
Variance reduction?	Manual	Manual	Manual	Built-in
Implementation	1 line	Lookup table	Formula	Platfo

Key Takeaway

Which method to use:

- **Closest to Eppo in these simulations:** in the conditions of this simulation, Pocock seems like the most reasonable choice among the alternative methods.
- **Avoid over-correction:** Bonferroni is often too conservative, reducing sensitivity more than needed.
- **Avoid early over-triggering:** O'Brien–Fleming is very conservative early, so early stopping is rare in this setup.

That said, we recommend running simulations, A/A tests, or analysing historical tests in each domain, using its specific circumstances (i.e. KPIs and their standard deviations) before deciding which alternative method to use in each domain.

Timing insight: the method choice matters most at the beginning of tests, when monitoring is mostly for implementation issues. As sample size grows, interval widths become more similar across methods.

Implementing the Hybrid Approach Without Eppo

Act 2 introduced the hybrid approach: sequential confidence interval on guardrail KPIs for early abort, standard confidence interval on the primary KPI at the planned end date. Below is how to implement it using any of the three correction methods above.

Step 1: Classify your metrics

Category	Examples	Testing method
Primary KPI	Conversion rate, sign-ups	Fixed-horizon ($z = 1.96$)
Guardrail KPIs	Revenue, error rate, latency, crash rate	Sequential (any of the 3 methods)

Step 2: Set guardrail significance levels

If you have G guardrail KPIs, set $\alpha_{\text{guardrail}} = \alpha / G$ for each one.

Example: 3 guardrails ($G = 3$), $\alpha = 0.05$, 4 weekly peeks ($K = 4$):

$$\alpha_{\text{per guardrail}} = \frac{0.05}{3} \approx 0.0167$$

Then apply Bonferroni within each guardrail:

$$z^* = \Phi^{-1}\left(1 - \tfrac{0.0167}{2\times 4}\right) = \Phi^{-1}(0.9979) \approx 2.86$$

Step 3: Run the experiment

At each scheduled peek:

1. For **each guardrail KPI**: compute the confidence interval using your chosen sequential method.
2. If **any** guardrail confidence interval is entirely on the harmful side of zero: **ABORT** the experiment.
3. Otherwise: continue.

At the **end** of the experiment:

1. For the **primary KPI**: compute a standard confidence interval with the full α:

$$\text{CI}_{\text{primary}} = \hat{\tau} \ \pm \ z_{\alpha/2} \cdot \text{SE}$$

SHIP if the confidence interval excludes zero. Otherwise: no significant effect.

Key Takeaway

The hybrid approach without Eppo:

- Apply Bonferroni across your G guardrail KPIs: $\alpha_g = \alpha/G$.
- Within each guardrail, use O'Brien–Fleming (best) or Bonferroni (simplest).
- Test the primary KPI once at the end with a standard confidence interval — no correction needed.
- **This is the recommended approach for teams without a sequential testing platform.**

Appendix

Show the code of the simulations

Summary

Act	Key message
1	Peeking inflates false positives. You need a test that is valid at every time you check.
2	Sequential confidence intervals are valid at every moment. The hybrid approach uses sequential confidence intervals for guardrails and classical confidence intervals for the primary KPI.
3	DIY alternatives: Bonferroni, Pocock, O'Brien–Fleming. All widen the confidence interval to control error when peeking.
4–6	Martingales and Ville's inequality: the math that makes peeking safe.
7–9	SPRT, mixture, and confidence sequences: the evolution of sequential testing and estimation.
10–13	Eppo's pipeline: regression adjustment, sequential confidence intervals, and actionable decisions.

The core trade-off: Valid interim monitoring requires wider confidence intervals. The penalty depends on the method, but the benefit is freedom to peek and act at any time.

Method	Price of peeking
No correction (naive)	Invalid inference (20–30% false positives)
Bonferroni	Confidence interval ~28–68% wider (depending on K)
Pocock	Confidence interval ~20–37% wider
O'Brien–Fleming	Confidence interval ~3% wider at final analysis

Method	Price of peeking
Eppo (sequential confidence interval)	Confidence interval ~10–40% wider, valid *continuously*
Hybrid (recommended)	No penalty on primary KPI; sequential penalty on guardrails only

Mathematical Reference

Notation

Symbol	Meaning	Where used
Y_i	Outcome for user i during the experiment	Act 2, Steps 2–3
X_i	Pre-experiment covariate for user i	Act 2, Step 3 (CUPED)
Y_i^*	Adjusted outcome after regression adjustment	Act 2, Step 3
$\hat{\tau}(t)$	Estimated treatment effect at time t	Act 2, Step 4
$\hat{\mu}_0(t),\ \hat{\mu}_1(t)$	Adjusted group means at time t	Act 2, Step 4
$\hat{\sigma}_{\hat{\tau}}(t)$	Estimated standard error of treatment effect	Act 2, Step 5
$s_g^2(t)$	Sample variance in group g at time t	Act 2, Step 5
$n_0(t),\ n_1(t)$	Number of users in control / treatment at time t	Throughout
n	Total observations: n_0(t) + n_1(t)	Act 2, Step 6
α	Significance level (typically 0.05)	Throughout

Symbol	Meaning	Where used
ν	Tuning parameter — calibrated to planned sample size	Act 2, Step 6 (sequential confidence interval)
M	Expected total sample size	Act 2, Step 6
K	Number of pre-planned analysis times (peeks)	Act 3
k	Current peek number (1 to K)	Act 3
t_k	Information fraction at peek k: k/K	Act 3 (OBF)
Z_k	z-statistic at peek k: estimated effect divided by its SE	Act 3
$z_{\alpha/2}$	Standard Normal critical value (≈1.96 for 95% CI)	Throughout
$z_{\alpha/(2K)}$	Bonferroni-adjusted critical value	Act 3 (Bonferroni)
c_P	Pocock constant critical value	Act 3 (Pocock)
c_k^{OBF}	O'Brien–Fleming critical value at peek k	Act 3 (OBF)
Φ^{-1}	Inverse standard Normal CDF (quantile function)	Act 3
ρ	Correlation between pre-experiment and experiment metrics	Act 2, Step 3
G	Number of guardrail KPIs	Act 3 (Hybrid)

Key Formulas

Standard Confidence Interval

$$\mathrm{CI} = \hat{\tau} \ \pm \ z_{\alpha/2} \cdot \mathrm{SE}$$

Valid only at one pre-specified time. The constant $z_{\alpha/2} \approx 1.96$ for a 95% confidence interval. A

confidence interval excluding zero is equivalent to rejecting the null hypothesis at level α.

Sequential Confidence Interval (Howard et al., 2021)

$$\mathrm{CI}(t) = \hat{\tau}(t) \;\pm\; \hat{\sigma}_{\hat{\tau}}(t) \cdot \sqrt{\frac{n+\nu}{n} \cdot \log\frac{n+\nu}{\nu\alpha}}$$

Valid at *all times simultaneously*. The sequential multiplier first decreases — reaching a minimum before n^* — then slowly rises. It always stays above 1.96.

CUPED Variance Reduction

$$Y_i^* = Y_i - \hat{f}_g(X_i) \qquad \mathrm{Var}(Y^*) = \mathrm{Var}(Y)(1-\rho^2)$$

At $\rho = 0.7$, $1 - 0.7^2 = 0.51$ — about 49% of the variance is eliminated.

Treatment Effect & Standard Error

$$\hat{\tau} = \bar{Y}_1 - \bar{Y}_0 \qquad \mathrm{SE} = \sqrt{\frac{s_0^2}{n_0} + \frac{s_1^2}{n_1}}$$

Bonferroni Adjusted Confidence Interval

$$\mathrm{CI}_k = \hat{\tau}(t_k) \;\pm\; z_{\alpha/(2K)} \cdot \mathrm{SE}(t_k)$$

O'Brien–Fleming Boundary

$$c_k^{\text{OBF}} = \frac{z_{\alpha/2}}{\sqrt{t_k}} = z_{\alpha/2} \cdot \sqrt{\frac{K}{k}}$$

Pocock Boundary

$$\mathbb{P}\Big(\max_{1 \le k \le K} |Z_k| > c_P\Big) = \alpha$$

With correlation: $\mathrm{Cor}(Z_j, Z_k) = \sqrt{\min(j,k)/\max(j,k)}$

Detailed Explanations

Confidence Sequences

A confidence sequence is a sequence of confidence sets $(C_t)_{t \ge 1}$ such that the true parameter lies in C_t for *all* t simultaneously with high probability. Unlike a standard confidence interval (valid at one time), a confidence sequence allows continuous monitoring. The price: wider intervals, especially early on.

Alpha Spending

The total false positive rate αcan be thought of as a “budget” that is spent across analyses. Bonferroni spends α/Kat each peek. O'Brien–Fleming spends almost nothing early and most at the end. The alpha-spending function $\alpha^*(t)$ formalises this: at information fraction t, cumulative type I error must not exceed $\alpha^*(t)$.

Information Fraction

The information fraction at peek k is $t_k = k/K$ for equally-spaced analyses, or more generally the ratio of current to planned total statistical information (proportional to sample size for equal variances). It ranges from 0 to 1 and determines where you are in the experiment.

FWER vs FDR

FWER (Family-Wise Error Rate): the probability of making *at least one* false positive across all tests. Bonferroni, Pocock, and OBF all control FWER. **FDR** (False Discovery Rate): the expected *proportion* of false positives among rejected hypotheses. FDR control is less strict and typically used when testing many hypotheses simultaneously (e.g. gene expression). For A/B testing with a small number of metrics, FWER control is standard.

References

1. **Deng, A., Xu, Y., Kohavi, R., & Walker, T.** (2013). Improving the sensitivity of online controlled experiments by utilizing pre-experiment data. *Proc. WSDM*, 123–132. — The original CUPED paper. Variance reduction via pre-experiment covariates.

2. **Howard, S.R., Ramdas, A., McAuliffe, J., & Sekhon, J.** (2021). Time-uniform, nonparametric, nonasymptotic confidence sequences. *Annals of Statistics*, 49(2), 1055–1080. — The confidence sequence framework that Eppo adopted.

3. **Pocock, S.J.** (1977). Group sequential methods in the design and analysis of clinical trials. *Biometrika*, 64(2), 191–199.

4. **O'Brien, P.C. & Fleming, T.R.** (1979). A multiple testing procedure for clinical trials. *Biometrics*, 35(3), 549–556.

5. **Schmit, M. & Miller, T.** (2024). Sequential confidence intervals for comparing two proportions with variance reduction. *Eppo Technical Report*. — Eppo's implementation

combining Howard et al.'s confidence sequences with generalised CUPED.

Comments

Sign in with GitHub to leave a comment. When commenting for the first time, please add your title and institution / workplace at the top of your message. Comments are moderated and appear after approval.

Contact us

Found an error? Have a suggestion for improvement? Please let us know!

Name

Email

Message

Send message